

Robustness Checks

Sarah Eckhardt

2025-01-31

0. Introduction

SIPP is representative at the national level, not sub-nationally. Given the sample size, 2+ dimensional cross-tabs may not be robust. Ideally, we would generate cuts by:

1. urban urban/rural vs industry
2. urban/rural vs states
3. generation
4. urban/rural by generation

In order to ascertain whether if any of these cuts are feasible, we conduct a series of robustness checks.

Robustness checks:

a. Simple sample size checks:

As a general rule, sub-groups should have no fewer than 30 obs (un-weighted).

b. Estimate sampling variability:

Compute standard errors for key estimates within each sub-group using survey weights. Large standard errors indicate a lack of reliability.

c. Design Effects

Check the design effect to see if clustering or stratification increases variance. High DEFF (>2) suggests that a larger sample may be needed for stability.

d. Margin of Error

Calculate confidence intervals for subgroup estimates. Wide confidence intervals indicate non-robustness.

1. Set-up

```

# remove dependencies
rm(list = ls())

# load packages
library(dplyr)
library(survey)
library(batman)
library(ggplot2)

# set user path
project_directories <- list(
  "sarah" = "/Users/sarah/Documents/GitHub/Retirement-Analysis-Urban-Rural"
)

current_user <- Sys.info()[["user"]]
if (!current_user %in% names(project_directories)) {
  stop("Root folder for current user is not defined.")
}

# set project paths
project_path = project_directories[[current_user]]
data_path = file.path(project_path, "Data")
output_path = file.path(project_path, "Output")

# Read in wrangled SIPP data (see 1. wrangle SIPP 2023.R)
load(paste(output_path, "SIPP_2023_WRANGLER.RData", sep = "/"))

# convert access, participation, and matching to bools
sipp_2023 = sipp_2023 %>%
  mutate(ANY_RETIREMENT_ACCESS_bool = case_when(
    ANY_RETIREMENT_ACCESS == "Yes" ~ 1,
    ANY_RETIREMENT_ACCESS == "No" ~ 0
  ),
  PARTICIPATING_bool = case_when(
    PARTICIPATING == "Yes" ~ 1,
    PARTICIPATING == "No" ~ 0
  ),
  MATCHING_bool = case_when(
    MATCHING == "Yes" ~ 1,
    MATCHING == "No" ~ 0
  ))

```

2. Robustness checks for urban/rural by major industry group

```

metro_ind_access = sipp_2023 %>% count(METRO_STATUS, INDUSTRY_BROAD,
                                       ANY_RETIREMENT_ACCESS) %>%

```

```

    filter(METRO_STATUS != "Not identified") %>% filter(!is.na(METRO_STATUS))

# hist(metro_ind_access$n, main = "metro X industry X access", xlab = "n")

# how many are under 30?
count(metro_ind_access %>% filter(n<30))

```

2.a Sample size check

```

##      n
## 1 14

```

```

# what do we exclude?
metro_ind_access_nu30 = metro_ind_access %>% filter(n<30)
unique(metro_ind_access_nu30$INDUSTRY_BROAD)

```

```

## [1] "Public Administration"
## [2] "Agriculture, Forestry, Fishing, and Hunting, and Mining"
## [3] "Arts, Entertainment, and Recreation, and Accommodation and Food Services"
## [4] "Finance and Insurance, and Real Estate and Rental and Leasing"
## [5] "Information"
## [6] "Other Services, Except Public Administration"
## [7] "Wholesale Trade"

```

```

excl = unique(metro_ind_access_nu30$INDUSTRY_BROAD)

# if we exclude unreliable industries, what does this leave us with?
metro_ind_access_n30 = metro_ind_access %>%
  filter(!(INDUSTRY_BROAD %in% excl))

unique(metro_ind_access_n30$INDUSTRY_BROAD)

```

```

## [1] "Construction"
## [2] "Educational Services, and Health Care and Social Assistance"
## [3] "Manufacturing"
## [4] "Retail Trade"
## [5] "Transportation and Warehousing, and Utilities"
## [6] "and Waste Management Services"

```

```

incl = unique(metro_ind_access_n30$INDUSTRY_BROAD)

```

```

# if we exclude these industries, is variance too large?
sipp_2023_excl_ind = sipp_2023 %>%
  filter(INDUSTRY_BROAD %in% incl) %>%
  filter(METRO_STATUS != "Not identified")

design <- svydesign(ids = ~0, weights = ~WPFINWGT, data = sipp_2023_excl_ind)

# mean, se of retirement access generally.
svymean(~ANY_RETIREMENT_ACCESS_bool, design)

```


2.b sampling variability

```
##                                mean      SE
## ANY_RETIREMENT_ACCESS_bool 0.59878 0.0077

# means for sub groups
res = svyby(~ANY_RETIREMENT_ACCESS_bool, by = ~ METRO_STATUS + INDUSTRY_BROAD, design, svymean)

# plot

ggplot(res, aes(x = interaction(METRO_STATUS, INDUSTRY_BROAD), y = ANY_RETIREMENT_ACCESS_bool, fill = a
  geom_bar(stat = "identity", position = position_dodge(), color = "black") +
  geom_errorbar(aes(ymin = ANY_RETIREMENT_ACCESS_bool - se, ymax = ANY_RETIREMENT_ACCESS_bool + se),
    width = 0.2, position = position_dodge(0.9)) +
  labs(title = "Retirement Access with Standard Errors",
    x = "Industry & Metro",
    y = "Retirement Access",
    fill = "Metro Status",
    axis.ticks.x = element_blank()) +
  theme_minimal()
```



```
# this looks ok, and fits w/ intuition; generally metro areas should have better access.
# SE for rural workers in information is quite high, as are point estimates. consider dropping.
```

```

# design effects --- does clustering or stratification increase variance?
# on the subsection of >30 n.
# DEFF ~ 1 similar to random sample
# > 1 higher variance w/ clustering
# >2 high clustering, less precise
# < 1 stratification improves precision

# consider: re-run svydesign construction to exclude Information

svymean(~ANY_RETIREMENT_ACCESS_bool, design, deff = TRUE)

```

2.c design effects

```

##                                mean      SE   DEff
## ANY_RETIREMENT_ACCESS_bool 0.5987818 0.0076924 1.4438

```


```

# all subgroups
res2 = svyby(~ANY_RETIREMENT_ACCESS_bool, by = ~METRO_STATUS + INDUSTRY_BROAD, design, svymean, deff = 'f')

# plot DEFF results
# this looks ok; DEFF < 2 for all sub-groups

ggplot(res2, aes(x = interaction(METRO_STATUS, INDUSTRY_BROAD), y = DEff.ANY_RETIREMENT_ACCESS_bool, fill = 'black')) +
  geom_bar(stat = "identity", position = position_dodge(), color = "black") +
  labs(title = "Design Effect (DEFF) by Industry and Metro",
       x = "Industry & Metro",
       y = "Design Effect (DEFF)",
       fill = "Metro Status") +
  theme_minimal()

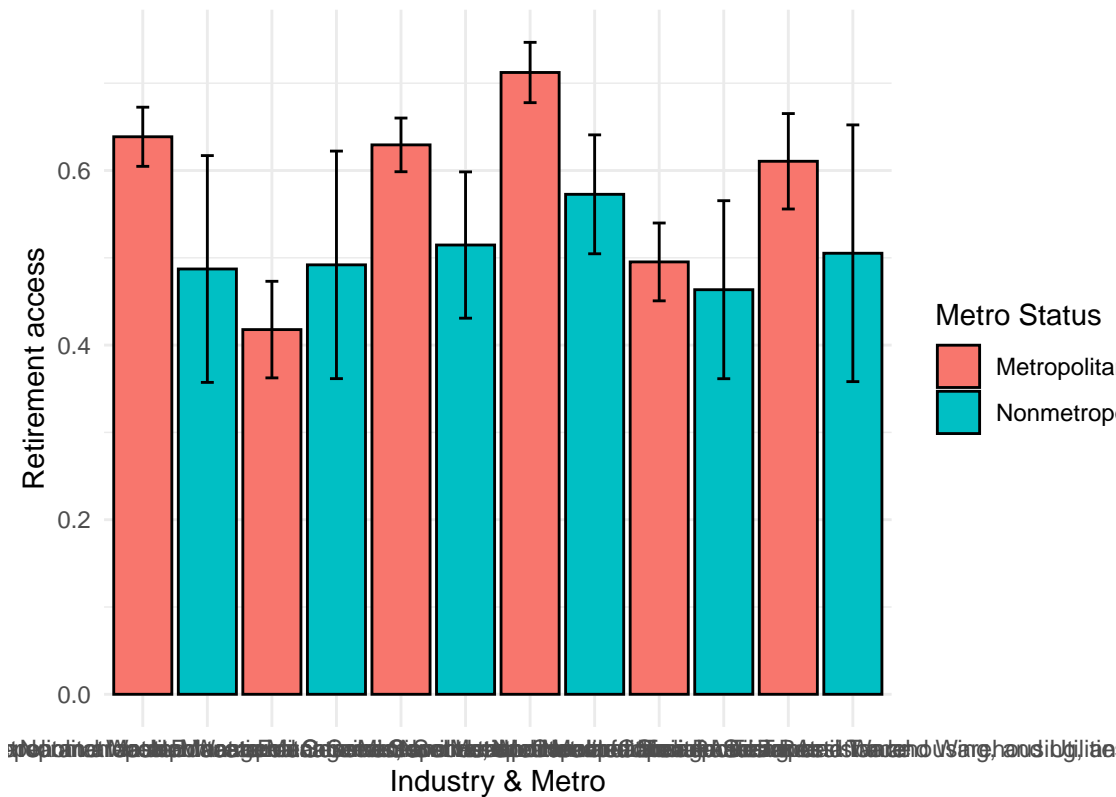
```



```
res2 = res2 %>%
  mutate(
    lower_CI = ANY_RETIREMENT_ACCESS_bool - 1.96 * se,
    upper_CI = ANY_RETIREMENT_ACCESS_bool + 1.96 * se
  )

ggplot(res2, aes(x = interaction(METRO_STATUS, INDUSTRY_BROAD), y = ANY_RETIREMENT_ACCESS_bool, fill = "Metro Status")) +
  geom_bar(stat = "identity", position = position_dodge(), color = "black") +
  geom_errorbar(aes(ymin = lower_CI, ymax = upper_CI),
    width = 0.2, position = position_dodge(0.9)) +
  labs(title = "Retirement access with 95% Confidence Intervals",
    x = "Industry & Metro",
    y = "Retirement access",
    fill = "Metro Status") +
  theme_minimal()
```

Retirement access with 95% Confidence Intervals



2.d confidence intervals

```
# confidence intervals for non-metro areas's information industry allow for estimates above 100% retirement access
confint(svymean(~ANY_RETIREMENT_ACCESS_bool, design))
```

```
##                2.5 %    97.5 %
## ANY_RETIREMENT_ACCESS_bool 0.5837049 0.6138587
```

```
# final list that passed robustness checks
```

```
print(incl)
```

```
## [1] "Construction"
## [2] "Educational Services, and Health Care and Social Assistance"
## [3] "Manufacturing"
## [4] "Retail Trade"
## [5] "Transportation and Warehousing, and Utilities"
## [6] "and Waste Management Services"
```

4. Robustness checks for States X Metro/non-metro

```
# generate a crosswalk of state abbrs, state fips
state_fips_lookup <- data.frame(
```

```


state_fips = c(1, 2, 4, 5, 6, 8, 9, 10, 11, 12, 13, 15, 16, 17, 18, 19,
               20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
               34, 35, 36, 37, 38, 39, 40, 41, 42, 44, 45, 46, 47, 48,
               49, 50, 51, 53, 54, 55, 56),
state_abbr = c("AL", "AK", "AZ", "AR", "CA", "CO", "CT", "DE", "DC", "FL", "GA",
               "HI", "ID", "IL", "IN", "IA", "KS", "KY", "LA", "ME", "MD", "MA",
               "MI", "MN", "MS", "MO", "MT", "NE", "NV", "NH", "NJ", "NM", "NY",
               "NC", "ND", "OH", "OK", "OR", "PA", "RI", "SC", "SD", "TN", "TX",
               "UT", "VT", "VA", "WA", "WV", "WI", "WY")
)

sipp_2023 = sipp_2023 %>% left_join(state_fips_lookup, by = c("TST_INTV" = "state_fips"))

# basic state count.
sipp_2023 %>%
  count(state_abbr, METRO_STATUS) %>%
  filter(!is.na(METRO_STATUS)) %>% filter(METRO_STATUS != "Not identified") %>%
  ggplot(aes(x = state_abbr, y = n, fill = as.factor(METRO_STATUS))) +
  geom_bar(stat = "identity", position = "dodge", color = "black") + theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  labs(title = "States X metro pairs unweighted count",
       x = "State", y = "n")

```


States X metro pairs unweighted count



2.a Sample size check


```
# how many are n<30 ?
```

```
sipp_2023 %>%
  count(state_abbr, METRO_STATUS) %>%
  filter(!is.na(METRO_STATUS)) %>% filter(METRO_STATUS != "Not identified") %>%
  filter(n<30) %>%
  ggplot(aes(x = state_abbr, y = n, fill = as.factor(METRO_STATUS))) +
  geom_bar(stat = "identity", position = "dodge", color = "black") + theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  labs(title = "States X metro pairs with n < 30",
       x = "State", y = "n")
```



```
# how many states does this impact?
```

```
st_met_n_under30 = sipp_2023 %>%
  count(state_abbr, METRO_STATUS) %>%
  filter(!is.na(METRO_STATUS)) %>% filter(METRO_STATUS != "Not identified") %>%
  filter(n<30)

length(unique(st_met_n_under30$state_abbr))
```

```
## [1] 27
```

```
# > 50% of the sample has states with either a metro or a non-metro sample of n < 30.
# this is not a viable sub group.
```